

PML - Course Project

December 26, 2015

Overview

The goal of this project is to use the Weight Lifting Exercise Dataset from to build a model that can predict if participants were performing a curl correctly or performing an incorrect motion in 1 of 5 ways. Information about the data can be found at <http://groupware.les.inf.puc-rio.br/har>. This write-up will discuss preparing the data for modeling, fitting the model, and predicting the 20 test cases.

Data Preparation

The initial data is in a raw form with some blanks and error codes. We set these to NA in our read statement.

```
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2

download.file(url='https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv',
              destfile='./PML_train.csv')

#Read in data setting empty and error values to NA.
WL <- read.csv('./PML_train.csv', na.strings = c("NA", "#DIV/0!", ""), as.is = TRUE)
WL <- data.frame(WL)
```

Create a training and test set to estimate out of sample error.

```
inTrain <- createDataPartition(y = WL$classe, p=.75, list = F)
WLtrain <- WL[inTrain,]
WLtest <- WL[-inTrain,]
```

The predictors are to be measurements from various sensors worn by the participants. Thus for prediction purposes I'll ignore the first 6 columns which are participant name, timestamps, and other experiment variables.

```
exp_vars <- grep("X|user_name|raw_timestamp_part_1|raw_timestamp_part_2|cvtd_timestamp|num_window", names(WLtrain))
WLtrain <- WLtrain[, -exp_vars]
```

A number of the predictors still have many missing values. There are only 221 complete observations all of the predictors. I've fit a Gradient Boosting Machine with the complete dataset with missing values, but accuracy was only 49.93% (not shown). Far greater accuracy can be achieved by removing the variables with missing values.

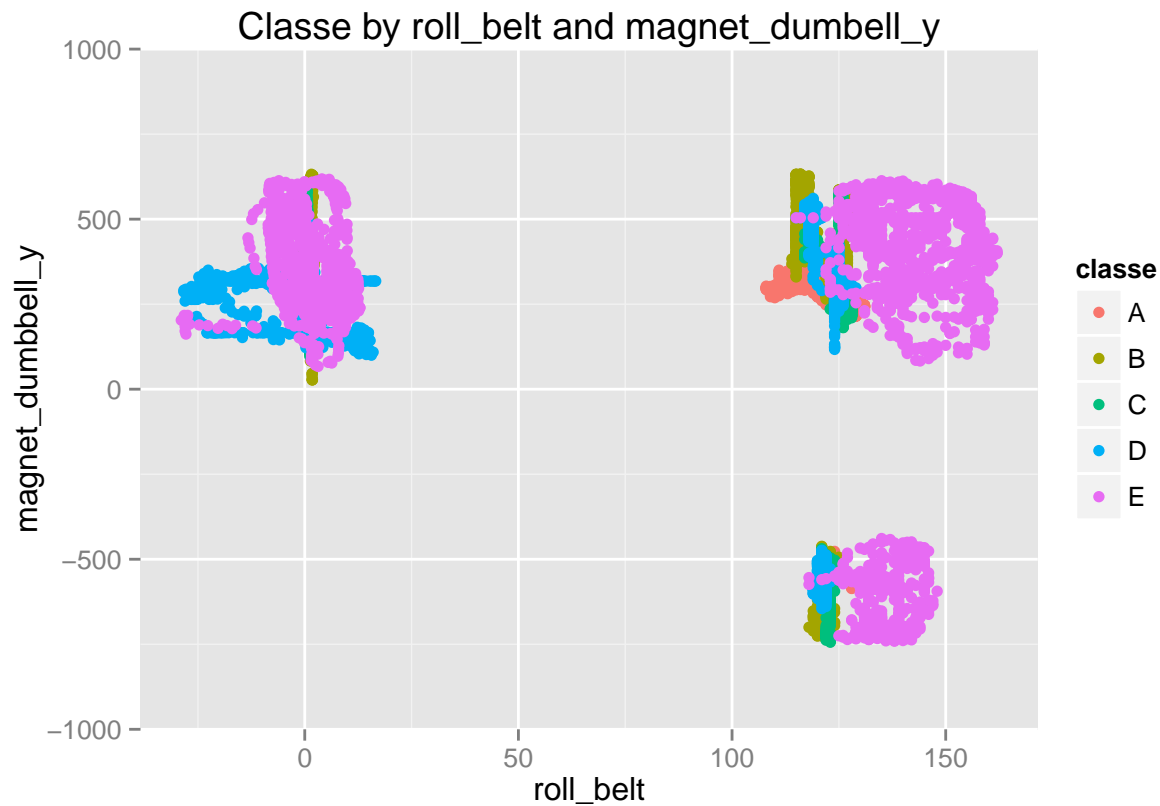
```
#clean variables with missing values
null_count <- sapply(WLtrain, function(x) sum(is.na(x)))
keep <- null_count==0
WLtrain <- WLtrain[, keep]
```

The prediction target is the classe variable. We'll make this a factor and look at a quick frequency and plot it by two of the most important variables.

```
WL$classe <- factor(WL$classe)
table(WL$classe)
```

```
##
##      A      B      C      D      E
## 5580 3797 3422 3216 3607
```

```
library(ggplot2)
ggplot(aes(x=roll_belt,y=magnet_dumbbell_y,colour=classe), data=WLtrain) + geom_point() + coord_cartesian(xlim=c(0,150),ylim=c(-1000,1000))
```



Modeling

We'll fit a random forest to the model. We'll use 5 fold cross validation in the train control to help mitigate potential overfitting.

```
Control1 <- trainControl(method = "cv", number = 5)
modFit1 <- train(classe ~ ., data=WLtrain, method = "rf", importance = T, trControl = Control1)
```


```
## Loading required package: randomForest
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
```

```
modFit1
```

```
## Random Forest
##
## 14718 samples
## 53 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 11774, 11775, 11773, 11775, 11775
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.9921864 0.9901157 0.001293926 0.001636516
## 27 0.9912354 0.9889124 0.001652847 0.002091711
## 53 0.9863434 0.9827231 0.002307971 0.002922301
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

We can examine the selection process as predictors were added.

```
plot(modFit1, log = "y", lwd = 2, main = "Random forest accuracy", xlab = "Predictors", ylab = "Accuracy")
```



The final model utilized 27 predictors. “Roll belt” and “magnet dumbbell y” were the most important predictors.

```
varImp(modFit1)
```

```
## rf variable importance
##
##   variables are sorted by maximum importance across the classes
##   only 20 most important variables shown (out of 53)
##
##           A      B      C      D      E
## roll_belt    81.81 100.00 95.00 93.74 75.81
## pitch_belt   80.71 94.49 76.39 79.93 81.12
## yaw_belt     89.22 93.54 85.92 91.85 65.96
## roll_arm     66.94 90.30 86.87 79.99 73.27
## magnet_dumbbell_z 83.60 85.11 89.72 77.81 74.15
## magnet_dumbbell_y 70.44 77.32 88.77 75.25 68.99
## accel_dumbbell_y 65.16 79.25 73.60 80.37 72.93
## pitch_forearm 66.15 80.15 77.80 77.10 70.06
## yaw_arm      64.33 78.03 67.43 71.69 66.66
## accel_dumbbell_z 65.99 77.69 71.29 69.28 69.74
## gyros_arm_y  57.60 77.32 63.47 68.36 59.00
## magnet_dumbbell_x 60.11 70.38 77.06 62.69 58.44
## gyros_forearm_y 55.56 76.87 64.26 70.14 61.72
## gyros_belt_z 63.77 76.33 70.89 73.08 69.72
## accel_belt_z 64.26 74.44 70.41 64.64 61.59
## magnet_forearm_z 60.18 74.41 68.65 72.02 66.32
## gyros_dumbbell_y 62.55 73.05 71.85 70.58 61.05
## roll_dumbbell 60.70 64.54 72.90 67.45 63.73
## gyros_dumbbell_x 61.35 72.33 68.00 60.77 59.18
## gyros_dumbbell_z 54.64 71.82 60.71 60.58 61.35
```

Out of Sample Error Prediction

To estimate the out of sample error we apply the prediction to the test set.

```
test.Predict = predict(modFit1, WLtest)
confusionMatrix(test.Predict, WLtest$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A      B      C      D      E
##           A 1395      5      0      0      0
##           B      0  937      4      0      0
##           C      0      7  850     10      0
##           D      0      0      1   794      0
##           E      0      0      0      0   901
##
## Overall Statistics
##
##           Accuracy : 0.9945
##           95% CI : (0.992, 0.9964)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##                      Kappa : 0.993
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  0.9874  0.9942  0.9876  1.0000
## Specificity      0.9986  0.9990  0.9958  0.9998  1.0000
## Pos Pred Value   0.9964  0.9957  0.9804  0.9987  1.0000
## Neg Pred Value   1.0000  0.9970  0.9988  0.9976  1.0000
## Prevalence       0.2845  0.1935  0.1743  0.1639  0.1837
## Detection Rate   0.2845  0.1911  0.1733  0.1619  0.1837
## Detection Prevalence 0.2855  0.1919  0.1768  0.1621  0.1837
## Balanced Accuracy 0.9993  0.9932  0.9950  0.9937  1.0000
```

The out of sample error prediction is 1-accuracy, which has an estimate of 0.86% with a 95% CI of (0.62%, 1.16%).